

With Claude Sonnet 5.5, a sharper prompt and one parser fix, the Policy Studio proposes the right rules for 90% of real NDAs it was never tuned on, up from 57%.

We ran 30 confidentiality agreements that companies filed with the SEC through the Policy Studio, 5 times each, and scored every run against an answer key of the rules a reviewer would expect. The changes were developed against the first 10. The other 20 were fetched afterwards in two batches and nothing was tuned on them, so their figure is the one to quote.

90.0% weighted accuracy on 20 agreements never used for tuning, 100 runs 95% range 75%–100%	90 / 100 runs with every rule exactly right	56.7% today's Studio on the same 20 agreements	9 s median time per agreement
---	---	--	---

How a run is scored. Each rule is identified by what it does on the agent: which tool, which folder or marked term, and which decision. A run's accuracy is correct rules ÷ (expected + generated – correct), so a missing rule, an extra rule and a rule with the wrong decision all count against it. Runs are weighted by how many rules their agreement should produce, and the 95% range comes from resampling agreements. These agreements mostly limit who may see the material, so most expect 2 rules: hold or block the agent reading and filing in the target's folder. One wrong decision makes a run score 0.

Four configurations, three sets of agreements

EXTRACTION	FIRST 10 changes developed here	SECOND 10 never tuned on	THIRD 10 never tuned on	SECOND + THIRD 95% range	EVERY RULE RIGHT second + third
Today: gpt-5-mini, current prompt	42.9%	56.0%	57.3%	56.7% (39–75%)	42 / 100
Claude Sonnet 5.5, current prompt	52.7%	70.0%	66.7%	68.3% (52–83%)	45 / 100
Claude Sonnet 5.5, new prompt	90.5%	90.0%	90.0%	90.0% (75–100%)	90 / 100
Claude Sonnet 5.5, new prompt with worked examples	100.0%	90.0%	87.3%	88.7% (74–100%)	86 / 100

- **Today's Studio** loses most clauses to one bug: gpt-5-mini returns curly apostrophes as a control character, so the check that every quoted value appears word for word in the NDA fails and the clause is dropped.
- **Claude Sonnet 5.5 alone** returns apostrophes intact and adds about 12 points.
- **The new prompt** adds another 22. It asks the model three questions the Studio used to answer with keyword lists: does consent unlock the action (however worded), is a breach of this clause itself a material breach, and does the clause really impose the duty it was classified as.

- **Worked examples in the prompt** made the first 10 perfect, the set they were written against, but did not help on fresh agreements: 87% on the third 10 against 90% without them. They fixed one kind of mixed clause and broke another. We recommend the prompt without them.
- **The parser fix** gives clauses unique ids when an amendment numbers its clauses from 1 again. Before it, the Studio read the amendment's clause 2 in place of the main confidentiality clause, and Care.com scored 0 in every run of every model.

Every run, Sonnet 5.5 with the new prompt

AGREEMENT target / counterparty	RUN 1	RUN 2	RUN 3	RUN 4	RUN 5	MEAN
FIRST 10: PROMPT WRITTEN AGAINST THESE						
H&E Equipment / United Rentals	100	100	100	100	100	100%
Inari Medical / Stryker	100	100	100	100	100	100%
Kronos Bio / Tang Capital	100	100	100	100	100	100%
PHX Minerals / WhiteHawk	0	0	0	0	0	0%
Potbelly / RaceTrac	100	100	100	100	100	100%
Profire Energy / CECO	100	100	100	100	100	100%
Servotronics / TransDigm (letter)	100	100	100	100	100	100%
Societal CDMO / QHP Capital	100	100	100	100	100	100%
TSR / BCforward	100	100	100	100	100	100%
Vector Group / JT International	100	100	100	100	100	100%
SECOND 10: FETCHED AFTER THE PROMPT WAS WRITTEN						
Care.com / IAC	100	100	100	100	100	100%
Collectors Universe / N. Turner	100	100	100	100	100	100%
Dunkin' Brands / Inspire Brands	0	0	0	0	0	0%
Forescout / Advent	100	100	100	100	100	100%
GMS / Home Depot	100	100	100	100	100	100%
Houghton Mifflin Harcourt / Veritas (letter)	100	100	100	100	100	100%
Michaels / Apollo (letter)	100	100	100	100	100	100%
Overseas Shipholding / Saltchuk	100	100	100	100	100	100%
Pluralsight / Vista Equity	100	100	100	100	100	100%
Spirit Airlines / JetBlue	100	100	100	100	100	100%

AGREEMENT target / counterparty	RUN 1	RUN 2	RUN 3	RUN 4	RUN 5	MEAN
THIRD 10: FETCHED LAST, USED FOR NOTHING BUT THIS TEST						
At Home / Hellman & Friedman	0	0	0	0	0	0%
Blue Apron / Wonder (letter)	100	100	100	100	100	100%
Castle Brands / Pernod Ricard	100	100	100	100	100	100%
Computer Task Group / Cegeka	100	100	100	100	100	100%
Huttig Building Products / Woodgrain	100	100	100	100	100	100%
KeyW / Jacobs	100	100	100	100	100	100%
Landsea Homes / New Home Co	100	100	100	100	100	100%
Papa Murphy's / MTY (letter)	100	100	100	100	100	100%
Tower International / Autokiniton	100	100	100	100	100	100%
Vocera / Stryker	100	100	100	100	100	100%

Each name links to the agreement on sec.gov: buyer-side M&A confidentiality agreements filed 2019–2025, 1,700 to 6,300 words, from retail, airlines, energy, software, healthcare, building products and more; five are letters with unnumbered paragraphs. 27 of the 30 come out right in all 5 runs, and the results barely vary between runs.

Where it still differs from the answer key

Two agreements differ from the key in every configuration. Read closely, both are judgement calls where the Studio's reading is defensible, so neither is changed in the code; both go to the lawyer reviewing the key.

1
agreement

Dunkin': is disclosure absolute, or can consent lift it?

The main sentence has no consent carve-out, which is why the key says BLOCK. But the same clause also says who may count as a Representative, and that consent governs who may receive the material, so the Studio's REQUIRE_APPROVAL has a basis.

"no portion of the Confidential Information will be disclosed by you to any other person ... except to your Representatives" · "your Representatives shall not, without the prior written consent of the Company, include any Person that is a franchisee, supplier or competitor"

1

agreement

KeyW: which of two overlapping duties sets the decision?

§2.1 allows disclosure beyond Representatives "as otherwise consented to in writing by the Company" (the key's REQUIRE_APPROVAL). §2.4 separately says the recipient will not disclose "to anyone, except to those ... Representatives", with no consent wording. The Studio applies the stricter of the two (BLOCK), which is the safe reading.

If the reviewer sides with the Studio on both, the shipped configuration proposed the right rules for every one of the 20 never-tuned agreements in all but one run. A prompt that asks which action each consent phrase covers was tried; it confirmed the Dunkin' reading rather than changing it, so it was not kept.

The verifier step still passes almost nothing at its 0.8 threshold, so every rule goes to a person for review, as the Studio's Apply screen expects.

What TypeSafe JEV adds

JEV already classifies every clause (91% agree with the key on whether a clause limits recipients). We tested two more uses of it on all 30 agreements.

USE OF JEV	WHAT IT DOES	RESULT
Recipient check	A yes/no per clause: does it limit who may receive the information, whatever else it does? Clauses it flags are read again as recipient limits; Claude drops the ones that are not.	Catches 12 of the 14 recipient limits the classifier filed under another kind. Across all 30: 92.8% (90.2% without). On the 20 never-tuned agreements: 90.3% (90.0% without): it fixes At Home and PHX and breaks KeyW, where a pure use clause then sets the decision.
Consent second opinion	JEV answers the consent question independently; rules where it and Claude disagree go to review.	Not useful: it flagged none of the 21 wrong rules and 26 correct ones. On its own JEV answers consent correctly for 88% of recipient clauses, Claude for 92%.

So JEV is the better tool for spotting what kind of duty a clause carries, and Claude for reading what the clause allows. The recipient check is worth keeping once a clause found this way is stopped from overriding the decision of the clause that grants access, which is the KeyW case.

TypeSafe JEV or the OpenAI Decisions API

The Studio asks a decision model to classify each clause, judge each tool's role and check each rule; Claude does the extraction. We swapped TypeSafe JEV for OpenAI's Decisions API (gpt-6-luna, public beta since 6 October 2026) and ran all 30 agreements again, 5 times each.

DECISION MODEL	EXTRACTION	ALL 30	NEVER TUNED ON (20)	EVERY RULE RIGHT	AGREEMENTS THAT VARY BETWEEN RUNS
TypeSafe JEV	Claude Sonnet 5.5	90.2%	90.0%	135 / 150	0 of 30
TypeSafe JEV + recipient check	Claude Sonnet 5.5	92.8%	90.3%	137 / 150	4 of 30
Decisions API, one pick per clause	Claude Sonnet 5.5	56.1%	50.0%	83 / 150	3 of 30

DECISION MODEL	EXTRACTION	ALL 30	NEVER TUNED ON (20)	EVERY RULE RIGHT	AGREEMENTS THAT VARY BETWEEN RUNS
Decisions API + recipient check	Claude Sonnet 5.5	93.2%	89.7%	139 / 150	1 of 30
Decisions API + recipient check	Claude Sonnet 5.5, prompt with worked examples	93.0%	89.3%	136 / 150	3 of 30
Decisions API + recipient check	GPT-6 Luna	87.5%	82.0%	131 / 150	3 of 30

- **Asked to pick one kind per clause, Luna files the main confidentiality clause under "use restriction"** in most agreements, because those clauses both limit use and say "disclose only to Representatives". A use restriction produces no rule, so every letter-form agreement and several others lost their rules.
- **Asked a yes/no instead, it is the best detector we tested.** "Does this clause limit who may receive the information?" agrees with the answer key on 95.3% of 532 clauses (JEV 92.9%, the one-pick classifier 91.4%). With that check added, the Decisions API matches JEV: 93.2% on all 30, 89.7% on the 20 never-tuned agreements.
- **Claude still reads consent better.** Luna answers "can consent lift this limit?" correctly for 86% of recipient clauses, Claude for 92%; as the extraction model Luna scores 6 points lower and varies more between runs.
- **Its probabilities are useful at the extremes.** Below 0.1 a recipient answer was wrong 1% of the time; above 0.9 it was right 80% of the time, so a high probability alone should not skip review.

The Decisions API costs \$0.10 per million input tokens with nothing charged for output; a whole agreement took about 9 seconds end to end with Claude extracting, as with JEV. The same two agreements fail with either decision model: Dunkin' (consent read into a clause where it covers copying) and KeyW (a use clause setting a stricter decision than the clause that grants access).

Live on staging

The final configuration (Claude Sonnet 5.5 with the new prompt, the OpenAI Decisions API with the recipient check, the parser and apply fixes) compiled all 30 agreements once and applied every proposed rule to a fresh agent on staging, the way the Studio's Apply button does: create the rule inactive, dry-run a case it must stop and a case it must allow, then switch it on.

61 / 61 rules created, dry-run and switched on Every rule passed both of its dry-runs, and reading the agent back shows all 61 active with the decision and conditions that were sent.	61 / 61 found unchanged when the same 30 agreements were applied again No duplicates and no errors: the Studio now reads every page of the agent's rules and saves three at a time.	28 / 30 agreements with exactly the right rules in this run The two misses are the known ones: Dunkin' and KeyW.
---	--	---

The rules are on the agent **policy-studio-final** in the staging dashboard under Agents → policy-studio-final → Authorize → Policies. Each rule is named after its agreement and clause, for example "NDA public3-castle §1 read_document 0063/96301/".

Production

On 9 October we registered an evaluation agent in the OpenBox test organisation on production (api.openbox.ai) and tried the same Apply test. Agent registration works; applying rules does not, for two reasons:

- **The policy-rule API is not on production yet.** `/agent/:id/policy-rule`, which the Studio uses to create, dry-run and activate each rule, answers "Cannot GET" on production: the route does not exist. It arrived on the backend's develop branch on 27 August (rule fragmentation); production runs main, which is 357 commits behind develop and has neither the route nor its migration. Staging runs develop, which is why the same test passes there.
- **The production key cannot manage the older policy endpoints either.** `/agent/:id/policies` and `/agent/:id/behavior-rule` answer 403, "Required permissions: read:agent_policy" / "read:agent_behavior_rule".

Compiling an NDA does not touch the platform, so accuracy is the same on any environment.

Applying on production needs the develop release (with the policy-rule API) to ship to production; the same script then runs it: `run_env.py --env .env.production`.

METHOD AND LIMITS

The agreements. Found with EDGAR full-text search: exhibits to Schedule TO and 14D-9 filings. The text is the filed exhibit with the HTML removed; where a filing sets a clause number in its own table cell, the number is joined to its text, as a PDF or Word copy would show it. One candidate was replaced because it turned out to be an exclusivity letter, not a confidentiality agreement. In the third set, 8 of the 10 agreements call for REQUIRE_APPROVAL and 2 for BLOCK, which is typical of buyer-side agreements.

What ran. The Policy Studio pipeline from branch `accuracy-fixes` (commit `fdafda3` plus fixes), not yet merged; only the extraction model and prompt differ between configurations, and the parser fix applies to every run on the third set and to the new-prompt runs on the first two. Classification and verification used TypeSafe JEV (`jev-1.13.0`). `gpt-5-mini` was reached through OpenRouter pinned to OpenAI; Claude through the Anthropic API with structured JSON output. Cost per agreement was not measured.

The answer key. Every clause was labelled with its kind of control and the decision its wording calls for, following a written guide, with the decisive words quoted for each label; the expected rules were derived from the labels by the Studio's own rule templates. One key was corrected during the work: Societal's clean-team amendment contains a flat access ban that the first labelling did not count. The key has not been reviewed by a lawyer; the hardest calls are BLOCK against REQUIRE_APPROVAL when a clause both forbids disclosure and allows it with consent.

Reproduce. `EVAL_EXTRACTOR=claude-sonnet-5-5 EXTRACT_PROMPT=v2 RUNS_DIR=... PYTHONPATH=<branch>/src run_eval.py --runs 5 --only ..., then summarize.py` with the same PYTHONPATH.